

output embedding size).

To predict step labels for masked-out steps at pre-training time, we consider two training objectives: (1) step classification, and (2) distribution matching. We describe them below in the context of Masked Step Modeling.

**Step classification loss.** We use the outputs of  $f_{VT}$  to represent an  $S$ -dimensional prediction distribution over steps, where  $S = |Y|$ . We form the target distribution by placing all probability mass on the best textual step description  $y_i^*$  for each clip  $v_i$  according to the weak supervision process. That is,

$$y_i^* = \operatorname{argmax}_{y \in Y} p(y | v_i). \quad (1)$$

We calculate the cross entropy between the predicted and target distributions for each masked out clip, yielding the following expression:

$$-\log([f_{VT}(V_{\setminus M})]_j) \quad (2)$$

where  $j$  is the index of  $y_i^*$  in  $Y$ , i.e., such that  $y_i^* = Y_j$ . To get the final training objective for a single masked video  $V_{\setminus M}$ , we sum over all indices  $i \in M$ , and minimize with respect to  $\theta$ .

**Distribution matching loss.** For this objective, we treat the distribution of step labels  $p(y_i | v_i)$  from weak supervision as the target distribution for each clip  $v_i$ . We then compute the KL Divergence between the prediction distribution  $f_{VT}(V_{\setminus M})$  and the target distribution  $p(y_i | v_i)$  as follows:

$$\sum_{j'=1}^S p(Y_{j'} | v_i) \log \frac{p(Y_{j'} | v_i)}{[f_{VT}(V_{\setminus M})]_{j'}} \quad (3)$$

We sum over all  $i \in M$  and minimize with respect to  $\theta$ . Following [13], we use only the top- $k$  steps in  $p(y_i | v_i)$  and set the probability of the remaining steps to 0.

Lin *et al.* [13] show that the distribution matching loss results in a slight improvement over step classification loss. For VideoTaskformer, we find both objectives to have similar performance and step classification outperforms distribution matching on some downstream tasks.

We use  $f_{VT}$  as a feature extractor (layer before softmax) to extract step representations for new video segments.

### 3.2. Downstream Tasks

To show that the step representations learned by VideoTaskformer capture task structure and semantics, we evaluate the representations on 6 downstream tasks—3 new tasks which we introduce (mistake step detection, mistake ordering detection, and long-term step forecasting) and 3 existing benchmarks (step classification, procedural activity recognition, and short-term step forecasting). We describe the dataset creation details for our 3 new benchmarks in Sec. 4.

**Mistake Detection.** A critical aspect of step representations that are successful at capturing the semantics and structure

of a task is that, from these representations, *correctness* of task execution can be verified. We consider two axes of correctness: content (what steps are portrayed in the video) and ordering (how the steps are temporally ordered). We introduce 2 new benchmark tasks to test these aspects of correctness.

• **Mistake step detection.** The goal of this task is to identify which step in a video is incorrect. More specifically, each input consists of a video  $V = \{v_1, \dots, v_K\}$  with  $K$  steps.  $V$  is identical to some unaltered video  $V_1$  that demonstrates a correctly executed task, except that step  $v_j$  (for some randomly selected  $j \in [1, \dots, K]$ ) is replaced with a random step from a different video  $V_2$ . The goal of the task is to predict the index  $j$  of the incorrect step in the video.

• **Mistake ordering detection.** In this task, the goal is to verify if the steps in a video are in the correct temporal order. The input consists of a video  $V = \{v_1, \dots, v_K\}$  with  $K$  steps. There is a 50% probability that  $V$  is identical to some (correctly ordered) video  $V_1 = \{v_1^1, \dots, v_K^1\}$ , and there is a 50% probability that the steps are randomly permuted. That is,  $v_i = v_{\pi_i}^1$  for some random permutation  $\pi$  of indices  $[1, \dots, K]$ . The goal of the task is to predict whether the steps are ordered correctly or are permuted.

**Step Forecasting.** As another way to evaluate how learned step representations capture task structure, we test the capabilities of our model in anticipating future steps given one or more clips of a video.

• **Short-term forecasting.** Consider a video  $V = \{v_1, \dots, v_n, v_{n+1}, \dots, v_K\}$  where  $v_i$  denotes a step, and  $V$  has step labels  $\{y_1, \dots, y_K\}$ , where  $y_i \in Y$ , the finite set of all step labels in the dataset. Short-term forecasting involves predicting the step label  $y_{n+1}$  given the previous  $n$  segments  $\{v_1, \dots, v_n\}$  [13].

• **Long-term step forecasting.** We introduce the challenging task of long-term step forecasting. Given a single step  $v_i$  in a video  $V = \{v_1, \dots, v_K\}$  with step labels  $\{y_1, \dots, y_K\}$ , the task is to predict the step labels for the next 5 steps, i.e.  $\{y_{i+1}, y_{i+2}, \dots, y_{i+5}\}$ . This task is particularly challenging since the network receives very little context—just a single step—and needs to leverage task information learned during training from watching multiple different ways of executing the same task.

**Procedural Activity Recognition.** The goal of this task is to recognize the procedural activity (i.e., task label) from a long instructional video. The input to the network is all the  $K$  video clips corresponding to the steps in a video,  $V = \{v_1, \dots, v_K\}$ . The task is to predict the video task label  $t \in \mathcal{T}$  where  $\mathcal{T}$  is the set of all task labels for all the videos in the dataset.

**Step Classification.** In this task, the goal is to predict the step label  $y_i \in Y$  given the video clip corresponding to step  $v_i$  from a video  $V = \{v_1, \dots, v_K\}$ . No context other than the single clip is given. Therefore, this task requires fine-